

Discussion: Identifiability and Optimization in Diagonal Kalman Delta Networks

July 29, 2026

1 Setup

The diagonal Kalman Delta Network (KDN) predicts a diagonal covariance and uses it to form an uncertainty-dependent delta-rule write. For key $\mathbf{k}_t$, transition $\boldsymbol{\alpha}_t$, process noise $\boldsymbol{\omega}_t$, and observation noise r_t ,

$$\hat{\mathbf{p}}_t = \boldsymbol{\alpha}_t^2 \odot \mathbf{p}_{t-1} + \boldsymbol{\omega}_t, \quad (1)$$

$$\boldsymbol{\kappa}_t = \frac{\hat{\mathbf{p}}_t \odot \mathbf{k}_t}{r_t + \sum_i \hat{p}_{t,i} k_{t,i}^2}. \quad (2)$$

The memory update remains a delta update,

$$\mathbf{S}_t = \mathbf{D}_t \mathbf{S}_{t-1} + \boldsymbol{\kappa}_t \left(\mathbf{v}_t - (\mathbf{D}_t \mathbf{S}_{t-1})^\top \mathbf{k}_t \right)^\top, \quad \mathbf{D}_t = \text{diag}(\boldsymbol{\alpha}_t). \quad (3)$$

The implementation additionally tracks diagonal information $\mathbf{c}_t = 1/\mathbf{p}_t$ with the calibrated projected update

$$\mathbf{c}_t = \frac{1}{\hat{\mathbf{p}}_t} + \frac{d_k}{r_t} \mathbf{k}_t^2. \quad (4)$$

This is a scan-compatible diagonal-information approximation, not an exact diagonal Kalman posterior.

2 The Scale Symmetry

For any constant $\lambda > 0$, consider

$$\mathbf{p}'_t = \lambda \mathbf{p}_t, \quad \boldsymbol{\omega}'_t = \lambda \boldsymbol{\omega}_t, \quad r'_t = \lambda r_t, \quad \mu' = \frac{\mu}{\lambda}, \quad (5)$$

where $\mu = \mathbf{c}_0$ is the initial information and hence $\mathbf{p}_0 = 1/\mu$. The predicted covariance scales as

$$\hat{\mathbf{p}}'_t = \boldsymbol{\alpha}_t^2 \odot (\lambda \mathbf{p}_{t-1}) + \lambda \boldsymbol{\omega}_t = \lambda \hat{\mathbf{p}}_t. \quad (6)$$

Consequently, the gain is invariant:

$$\boldsymbol{\kappa}'_t = \frac{\lambda \hat{\mathbf{p}}_t \odot \mathbf{k}_t}{\lambda r_t + \sum_i \lambda \hat{p}_{t,i} k_{t,i}^2} = \boldsymbol{\kappa}_t. \quad (7)$$

The information update also transforms consistently as $\mathbf{c}'_t = \mathbf{c}_t/\lambda$. Therefore $\boldsymbol{\omega}$, r , and μ contain a common covariance-scale degree of freedom that cannot be identified from the memory output. Positive floors weaken the symmetry near their boundaries but do not remove the general conditioning problem.

The transition $\boldsymbol{\alpha}_t$ is different. It affects both covariance prediction and the memory transition $\mathbf{D}_t \mathbf{S}_{t-1}$, so it is not part of the exact scale symmetry. Nevertheless, $\boldsymbol{\alpha}_t$ and $\boldsymbol{\omega}_t$ are partially confounded in the covariance path because multiple pairs can produce similar $\hat{\mathbf{p}}_t$.

3 What the Noisy-MQAR Results Establish

The original diagonal model obtains 76.2 ± 22.3 percent over the 5×5 initialization and data grid. Its initialization component of variation (19.2 points) is larger than its data component (11.4 points). This is not uniform under-capacity: many cells reach 88–96 percent, while a few fail severely.

The most informative intervention is isotropic gain initialization. Zeroing the output map of the process-noise projection makes $\boldsymbol{\omega}_t$ uniform across channels at initialization, after which anisotropy is learned. At the same learning rate and with the same optimizer, this changes performance from

$$76.2 \pm 22.3 \quad \longrightarrow \quad 83.7 \pm 10.4, \quad (8)$$

and halves the initialization variation from 19.2 to 9.0 points. It also rescues both severe collapse cells. Thus the current evidence identifies random initial gain anisotropy as the primary failure, not the nominal learning rate or scale symmetry.

The scale symmetry remains relevant as a secondary optimization defect. It creates a flat direction, gives equivalent physical models different raw parameter values, and lets coupled quantities evolve under different optimizer geometries. In the current Muon setup, matrix weights use Muon with adjusted learning rates, while biases and the one-dimensional prior parameter use AdamW. A single printed learning rate therefore does not imply uniform motion of the physical uncertainty variables.

The prior μ is unlikely to explain late failures. It only controls the initial information, and its influence should decay rapidly over a long sequence when process noise has a positive floor. Learning it adds a weakly identified parameter with little expected benefit.

4 Ways to Remove or Reduce the Symmetry

4.1 Fix the observation-noise scale

The strongest and simplest gauge choice is

$$r_t \equiv 1, \quad \mu \equiv 1. \quad (9)$$

The model then learns process uncertainty in units of observation noise:

$$\boldsymbol{\omega}_t = \boldsymbol{\omega}_{\min} + \text{softplus}(f_{\omega}(\mathbf{x}_t)). \quad (10)$$

Because r can no longer be rescaled, the common covariance-scale symmetry is removed. This parameterization is the preferred baseline unless token-dependent observation reliability is shown empirically to matter.

4.2 Learn bounded relative observation reliability

If heteroscedastic observation noise is useful, it should change confidence relative to a fixed reference rather than learn an unrestricted global scale. One option is

$$r_t = \exp(s \tanh f_r(\mathbf{x}_t)), \quad (11)$$

with fixed s . For $s = \log 4$, $r_t \in [1/4, 4]$. The projection should have no learnable global bias, and the prior can remain fixed at $\mu = 1$. A soft anchoring penalty can further control drift,

$$\mathcal{L}_{\text{scale}} = \gamma (\mathbb{E}[\log r_t])^2. \quad (12)$$

This reduces rather than mathematically eliminates every possible input-dependent rescaling, but it prevents the practically harmful unbounded common-scale direction.

4.3 Parameterize a stationary uncertainty level

The α - ω interaction can be made more interpretable with

$$\omega_{t,i} = \omega_{\min} + (1 - \alpha_{t,i}^2)\rho_{t,i}, \quad \rho_{t,i} = \text{softplus}(f_\rho(\mathbf{x}_t)). \quad (13)$$

Here α controls retention and ρ is the uncertainty level toward which prediction moves. This is exactly the completed **DiagKLA1** experiment. Coupling alone changes 76.2 ± 22.3 to 76.9 ± 20.9 , which is effectively a wash. The completed **DiagKLA1-zi** experiment combines coupling with isotropic gain initialization and obtains 80.1 ± 13.3 , below the simpler **DiagKLA-zi** result of 83.7 ± 10.4 . Thus this parameterization has already been tested: it neither replaces isotropic initialization nor improves on free additive process noise once that initialization is used.

5 Recommended Ablation

The transition-coupled process-noise comparison is already complete:

Variant	Free additive ω	$(1 - \alpha^2)\rho$
Random anisotropic initialization	76.2 ± 22.3	76.9 ± 20.9
Isotropic gain initialization	83.7 ± 10.4	80.1 ± 13.3

The remaining experiment should isolate covariance-scale choices while keeping free additive ω and isotropic gain initialization fixed:

Variant	μ	r_t	ω_t
A	learned	learned positive	free additive
B	fixed at 1	learned positive	free additive
C	fixed at 1	fixed at 1	free additive
D	fixed at 1	bounded relative	free additive

Run at least the same 5×5 initialization–data grid. Report mean, total standard deviation, initialization variation, non-finite runs, and convergence time. During training, $\log \omega$, r , effective write strength, gain anisotropy, the minimum channel gain, and gradient norms for the process-noise, observation-noise, prior, and transition parameters.

If C matches or improves A, the production model should remove both r_{proj} and μ_{param} . If D improves C, retain only bounded relative observation reliability. Only after fixing this gauge is a separate uncertainty-parameter learning-rate sweep readily interpretable.

6 Why DiagKLA Does Not Yet Dominate IsoKLA

The available implementation tests establish that the Triton kernels reproduce the intended PyTorch recurrence; they do not establish that this recurrence is the best approximation for an associative memory. Empirically, the extra diagonal uncertainty state has not produced a general advantage over the isotropic state. This section separates the approximation question from the kernel-correctness question and states the experiment needed to identify the useful part of the diagonal model.

6.1 Current evidence

The controlled 220M/20B information-scale runs do not establish a general DiagKLA advantage. The best default-calibrated variants give:

Metric	DiagKLA, $s = d_k$	IsoKLA, $s = 1$
WikiText word perplexity ↓	28.99	28.79
LAMBADA perplexity ↓	35.12	35.66
Five-task multiple-choice average ↑	0.472	0.477
S-NIAH-2 at 4K ↑	0.272	0.046
S-NIAH-3 at 2K ↑	0.004	0.008
MK-NIAH-1 at 4K ↑	0.178	0.182

General language-modeling quality is mixed and effectively tied. DiagKLA has a large, specific S-NIAH-2 advantage, a direction that also appears in the earlier 220M comparison, but it does not improve the harder single-key or multikey probes. The recent sweeps use one training seed and different GPU world sizes, and a current-code 750M DiagKLA evaluation is not yet available. The evidence therefore supports a task-specific single-key retrieval signal, not a general DiagKLA ranking above IsoKLA.

The noisy-MQAR grid gives a similar result. Uniform process-noise initialization raises DiagKLA from 76.2 ± 22.3 to 83.7 ± 10.4 , while the best stabilized IsoKLA variant obtains 83.0 ± 8.4 . These are effectively tied at this resolution. The older trained-gain probe that reports large off-key gain and over-writing predates the current initialization and memory kernel; it motivates the mechanism test below but should not be treated as a diagnosis of the current default.

6.2 Approximation and calibration

With predicted diagonal covariance $\hat{\mathbf{P}}_t = \text{diag}(\hat{\mathbf{p}}_t)$, the exact Gaussian measurement update produces a dense posterior whose diagonal is

$$p_{t,i}^{\text{exact}} = \hat{p}_{t,i} - \frac{\hat{p}_{t,i}^2 k_{t,i}^2}{r_t + \sum_j \hat{p}_{t,j} k_{t,j}^2}. \quad (14)$$

The scan-compatible DiagKLA state instead uses

$$p_{t,i}^{\text{proj}} = \left(\hat{p}_{t,i}^{-1} + s \frac{k_{t,i}^2}{r_t} \right)^{-1}, \quad s = d_k \text{ by default.} \quad (15)$$

Even at $s = 1$, Equation (15) is not the diagonal of the exact posterior in Equation (14). For $s \neq 1$, the gain and information update additionally correspond to different measurement scales: the

current gain retains precision $1/r_t$, while the state used by future gains receives precision s/r_t . Thus $s = d_k$ is a dimension calibration of a gain-conditioning state, not an exact Kalman consequence.

For a spread-out unit key, $\mathbb{E}[k_i^2] = 1/d_k$, so $s = d_k$ gives a unit expected per-coordinate increment. Averaging these calibrated coordinates recovers the IsoKLA scale,

$$\frac{1}{d_k} \sum_i d_k \frac{k_{t,i}^2}{r_t} = \frac{\|\mathbf{k}_t\|_2^2}{r_t} = \frac{1}{r_t}. \quad (16)$$

Matching mean information does not match covariance or gain because inversion is nonlinear. Moreover, for a uniformly distributed unit key,

$$\text{Var}(d_k k_i^2) = \frac{2(d_k - 1)}{d_k + 2} \longrightarrow 2. \quad (17)$$

The coordinate noise therefore does not vanish with width. DiagKLA can interpret coordinate energy fluctuations as persistent uncertainty structure. The projection is also basis dependent: $(1, 1)/\sqrt{2}$ and $(1, -1)/\sqrt{2}$ are orthogonal but have identical elementwise squares and hence identical information increments. Whether the learned key basis aligns this diagonal statistic with useful uncertainty is an empirical question.

6.3 Information scale separates present writing from future protection

The role of s is clearest after separating the memory update at token t from the uncertainty state passed to later tokens. Define

$$z_t = \sum_i \hat{p}_{t,i} k_{t,i}^2, \quad \beta_t^{\text{eff}} = \mathbf{k}_t^\top \mathbf{k}_t = \frac{z_t}{r_t + z_t}. \quad (18)$$

Writing $\tilde{\mathbf{S}}_t = \mathbf{D}_t \mathbf{S}_{t-1}$ and $\mathbf{e}_t = \mathbf{v}_t - \tilde{\mathbf{S}}_t^\top \mathbf{k}_t$, the memory recurrence implies

$$\mathbf{S}_t = \tilde{\mathbf{S}}_t + \mathbf{k}_t \mathbf{e}_t^\top, \quad \mathbf{v}_t - \mathbf{S}_t^\top \mathbf{k}_t = (1 - \beta_t^{\text{eff}}) \mathbf{e}_t. \quad (19)$$

Thus β_t^{eff} is the fraction of the current residual corrected at its own key. The same $\mathbf{k}_t$ controls both the value write $+\mathbf{k}_t \mathbf{v}_t^\top$ and the erasure $-\mathbf{k}_t \mathbf{k}_t^\top \tilde{\mathbf{S}}_t$.

To expose causal ordering, temporarily allow a token-specific scale s_t . Current DiagKLA first computes $\hat{\mathbf{p}}_t$ and $\mathbf{k}_t$ from the incoming state $\mathbf{c}_{t-1}$, and only afterward computes the tracked projected covariance below. In this subsection, plain $\mathbf{p}_t$ denotes that DiagKLA state, not the exact posterior diagonal.

$$p_{t,i}(s_t) = \left(\hat{p}_{t,i}^{-1} + s_t \frac{k_{t,i}^2}{r_t} \right)^{-1} = \frac{\hat{p}_{t,i}}{1 + s_t \hat{p}_{t,i} k_{t,i}^2 / r_t}. \quad (20)$$

Consequently, for a local intervention that holds the incoming state and token projections fixed,

$$\frac{\partial \mathbf{k}_t}{\partial s_t} = 0, \quad \frac{\partial \mathbf{S}_t}{\partial s_t} = 0, \quad \frac{\partial p_{t,i}}{\partial s_t} = -\frac{k_{t,i}^2}{r_t} p_{t,i}^2 \leq 0. \quad (21)$$

The scale at token t therefore does not write that observation harder or change its same-token residual. It only increases the confidence state supplied to the future. In causal order, $\mathbf{c}_{t-1}$ determines $\hat{\mathbf{p}}_t$, which determines $\mathbf{k}_t$ and $\mathbf{S}_t$; s_t then enters $\mathbf{c}_t$ and can first affect the memory gain at token $t + 1$.

This token-local statement must not be confused with a comparison between two complete runs that use different global values of s . Except at the first token, a global s has already changed $\mathbf{c}_{t-1}$

and therefore generally changes the current gain. Conditional on fixed projected $\mathbf{k}$, α , ω , and r , its first delayed effect is explicit:

$$\frac{\partial z_{t+1}}{\partial s_t} = -\frac{1}{r_t} \sum_i \alpha_{t+1,i}^2 p_{t,i}^2 k_{t,i}^2 k_{t+1,i}^2 \leq 0, \quad (22)$$

$$\frac{\partial \beta_{t+1}^{\text{eff}}}{\partial s_t} = \frac{r_{t+1}}{(r_{t+1} + z_{t+1})^2} \frac{\partial z_{t+1}}{\partial s_t} \leq 0. \quad (23)$$

More generally, for $u > t$ the channelwise sensitivity propagates as

$$\frac{\partial p_{u,i}}{\partial s_t} = -\frac{k_{t,i}^2}{r_t} p_{t,i}^2 \prod_{n=t+1}^u \frac{\alpha_{n,i}^2}{(1 + s_n \hat{p}_{n,i} k_{n,i}^2 / r_n)^2}. \quad (24)$$

A larger past scale therefore reduces the later effective correction $\mathbf{k}_u^\top \boldsymbol{\kappa}_u$. It need not reduce every component or the norm of an anisotropic $\boldsymbol{\kappa}_u$: reducing the numerator coordinates also reduces their shared denominator. End-to-end trained runs need not be monotone either, because their learned projections can compensate for the changed recurrence.

Equations (23)–(24) also show what is being protected. Suppression is controlled by elementwise-energy overlap such as $\sum_i p_{t,i}^2 k_{t,i}^2 k_{u,i}^2$, not by the signed directional overlap $(\mathbf{k}_t^\top \mathbf{k}_u)^2$. At fixed projections, a future key with disjoint coordinate support receives no direct effective-gain suppression through this pathway, although the full memory can still differ through intervening writes. Dense orthogonal keys can instead have large energy overlap, including the maximum compatible with orthogonality in the two-dimensional example above. The tracker gives those two keys identical confidence increments and, under a symmetric incoming covariance, identical effective-gain suppression. Their signed memory updates and residuals need not be identical. Large s can consequently suppress a genuinely new orthogonal association rather than selectively preserve only the old one.

Prediction limits the duration of this protection. As $s_t \rightarrow \infty$, $p_{t,i} \rightarrow 0$ wherever $k_{t,i} \neq 0$, but

$$\hat{p}_{t+1,i} = \alpha_{t+1,i}^2 p_{t,i} + \omega_{t+1,i} \longrightarrow \omega_{t+1,i}. \quad (25)$$

Positive process noise restores a floor on uncertainty and hence future plasticity; small α also rapidly attenuates the confidence sensitivity in Equation (24) while directly decaying the memory. Thus larger s does not store more content or create capacity by itself. It changes the retention–plasticity operating point until ω and α wash out that change. If $\omega = 0$ and $\alpha \simeq 1$, an extreme s can instead freeze future gains on the affected coordinates.

The different optima for DiagKLA and IsoKLA are consistent with their different recurrences. For a spread-out normalized key, the DiagKLA increment in one coordinate has expectation $s/(d_k r_t)$, so $s = d_k$ cancels the coordinate dilution without changing the current DiagKLA gain locally. IsoKLA has one scalar state and $\|\mathbf{k}_t\|_2^2 = 1$, so its information increment is already s/r_t with no $1/d_k$ dilution. Moreover, the current IsoKLA implementation uses, locally holding the incoming $\hat{b}_t$ fixed,

$$\beta_t^{\text{iso}}(s) = \frac{\hat{b}_t}{r_t + s \hat{b}_t}, \quad c_t^{\text{iso}}(s) = \hat{b}_t^{-1} + \frac{s}{r_t}, \quad \left. \frac{\partial \beta_t^{\text{iso}}}{\partial s} \right|_{\hat{b}_t} = -\frac{\hat{b}_t^2}{(r_t + s \hat{b}_t)^2} < 0. \quad (26)$$

A token-local increase in s in IsoKLA therefore weakens the current write and increases future confidence at the same time. With fixed projections, the lower uncertainty inherited from past larger scales reinforces this reduction in the global- s recurrence. At $s = d_k$ these two anti-plasticity effects can be excessive, whereas DiagKLA obtains delayed protection without the same local write

penalty. This gives a mechanistic interpretation of the observed optima $s = d_k$ and $s = 1$, but it is not yet causal evidence that retention alone produces the downstream improvements.

A coherent Gaussian observation-precision multiplier has a different meaning. For a local intervention that holds the predicted prior $\hat{\mathbf{P}}_t$ fixed, increasing precision by s means using $r_t^{\text{eff}} = r_t/s$ in both the current gain and posterior. With $z_t = \mathbf{k}_t^\top \hat{\mathbf{P}}_t \mathbf{k}_t$, this gives

$$\boldsymbol{\kappa}_t^{\text{coh}}(s) = \frac{s\hat{\mathbf{P}}_t \mathbf{k}_t}{r_t + sz_t}, \quad \beta_t^{\text{coh}}(s) = \frac{sz_t}{r_t + sz_t}, \quad \frac{\partial \beta_t^{\text{coh}}}{\partial s} = \frac{r_t z_t}{(r_t + sz_t)^2} > 0. \quad (27)$$

The remaining current residual is then $1 - \beta_t^{\text{coh}} = r_t/(r_t + sz_t)$, so a more precise observation is assimilated more aggressively. At the same time, its exact posterior is

$$\mathbf{P}_t(s) = \hat{\mathbf{P}}_t - \frac{s\hat{\mathbf{P}}_t \mathbf{k}_t \mathbf{k}_t^\top \hat{\mathbf{P}}_t}{r_t + sz_t}, \quad \frac{\partial \ell^\top \mathbf{P}_t(s) \ell}{\partial s} = -\frac{r_t (\ell^\top \hat{\mathbf{P}}_t \mathbf{k}_t)^2}{(r_t + sz_t)^2} \leq 0. \quad (28)$$

Thus coherent higher precision has the two temporally distinct effects expected of a Kalman filter: write the observation harder now, then reduce later corrections in covariance-overlapping directions. Neither the current DiagKLA split nor Equation (26) for $s \neq 1$ has this interpretation. A coherent implementation should use r_t/s once; replacing r_t by r_t/s and also retaining a separate factor s in the information update would double-count the scale and produce an s^2 increment.

These mechanisms can be separated explicitly with a current-write scale s_{write} and a post-write confidence scale s_{protect} :

$$\boldsymbol{\kappa}_t(s_{\text{write}}) = \frac{s_{\text{write}} \hat{\mathbf{P}}_t \odot \mathbf{k}_t}{r_t + s_{\text{write}} z_t}, \quad c_{t,i} = \hat{p}_{t,i}^{-1} + s_{\text{protect}} \frac{k_{t,i}^2}{r_t}. \quad (29)$$

Current DiagKLA is $(s_{\text{write}}, s_{\text{protect}}) = (1, s)$. Setting both to s makes the scale consistent within the diagonal precision projection, although that projection still is not the exact posterior diagonal. A factorial comparison of $(1, 1)$, $(1, d_k)$, $(d_k, 1)$, and (d_k, d_k) would separate current assimilation from future protection; because $s_{\text{write}} = d_k$ may saturate the current correction, a gain-matched intermediate value should also be included. Log the same-token β_t^{eff} and residual contraction, the next-token and repeated-key gains, write/erase magnitudes, covariance quantiles, and effective memory half-life. Repeated-key, replacement, and orthogonal-key tasks can then identify whether the benefit comes from retention, impaired rewriting, or diagonal geometry.

A large learned r_t could make the $s = 1$ confidence increment too weak, but the preferred $s = d_k$ does not by itself diagnose observation noise as too large: increasing r_t also weakens the current gain, whereas increasing only s_{protect} does not. The relevant dimensionless quantities are z_t/r_t for the current write and $s\hat{p}_{t,i}k_{t,i}^2/r_t$ for the confidence update. These, together with $\omega_{t,i}/\hat{p}_{t,i}$, should be measured before attributing the result to r_t . Likewise, if r_t is freely learned, a globally coherent s is largely reparameterizable into r_t/s ; the present decoupled scale is not redundant because it changes protection without changing that observation’s same-token write locally.

6.4 An exact posterior-diagonal alternative

Although the full posterior is dense, its diagonal after one measurement can be computed without materializing a $d_k \times d_k$ covariance. Define

$$D_t = r_t + \sum_j \hat{p}_{t,j} k_{t,j}^2, \quad q_{t,i} = \hat{p}_{t,i} k_{t,i}^2. \quad (30)$$

The denominator D_t is already required for the diagonal Kalman gain. The exact posterior diagonal from Equation (14) can therefore be updated as

$$p_{t,i}^{\text{mom}} = \hat{p}_{t,i} \frac{D_t - q_{t,i}}{D_t} = \hat{p}_{t,i} (1 - \kappa_{t,i} k_{t,i}), \quad \kappa_{t,i} = \frac{\hat{p}_{t,i} k_{t,i}}{D_t}. \quad (31)$$

This costs $O(d_k)$ state and arithmetic per token. The multiplicative form in Equation (31) is preferable to subtracting two nearly equal covariance terms. Since $D_t - q_{t,i} = r_t + \sum_{j \neq i} q_{t,j} > 0$, it also makes positivity explicit. Equivalently, the marginal information update is

$$\frac{1}{p_{t,i}^{\text{mom}}} = \frac{1}{\hat{p}_{t,i}} + \frac{k_{t,i}^2}{r_t + \sum_{j \neq i} \hat{p}_{t,j} k_{t,j}^2}. \quad (32)$$

Unlike the current update, this expression retains the competition from the other key coordinates instead of replacing its denominator by r_t .

The word “exact” applies only to the diagonal of this one posterior when the predicted prior is diagonal. The same update creates off-diagonal entries

$$P_{t,ij} = -\frac{\hat{p}_{t,i} \hat{p}_{t,j} k_{t,i} k_{t,j}}{D_t}, \quad i \neq j. \quad (33)$$

These entries affect the next gain. Discarding them and carrying only $\mathbf{p}_t^{\text{mom}}$ is therefore a moment-projected, diagonal assumed-density filter, not the exact multi-token Kalman filter. Exact filtering across tokens requires the dense $\mathbf{P}_t$, with $O(d_k^2)$ state and arithmetic per token, as in ExactKLA.

This alternative also changes the parallelization structure. All channels in Equation (31) are coupled through D_t , so the recurrence does not decompose into the independent 2×2 Möbius maps used by the current gain scan. The straightforward implementation is a token-sequential $O(Td_k)$ gain recurrence, preferably fused in Triton; after the gains are produced, the existing chunked Kalman memory kernel remains applicable. Thus the exact posterior-diagonal update saves the dense $O(d_k^2)$ covariance but trades away the current logarithmic-depth prefix scan for the gain.

For the moment projection, the coherent scale from Subsection 6.3 specializes to $r_t^{\text{eff}} = r_t/s$ and gives

$$D_t^{(s)} = \frac{r_t}{s} + \sum_j \hat{p}_{t,j} k_{t,j}^2, \quad \kappa_{t,i}^{(s)} = \frac{\hat{p}_{t,i} k_{t,i}}{D_t^{(s)}}, \quad p_{t,i}^{(s)} = \hat{p}_{t,i} - \frac{\hat{p}_{t,i}^2 k_{t,i}^2}{D_t^{(s)}}. \quad (34)$$

This is the exact-posterior-diagonal specialization of Equations (27)–(28). A direct ablation should compare the current Möbius precision scan against this recurrent moment projection first with $s = 1$, and then, separately, with the coherent calibrated choice $r_t^{\text{eff}} = r_t/d_k$.

Is the moment projection likely to be a better mixer? It is a better approximation to the one-step posterior covariance, but it is not necessarily a better drop-in strategy for associative memory. Consider the symmetric case $\hat{p}_{t,i} = \hat{p}_t$ and $k_{t,i}^2 = 1/d_k$, and define the effective write strength

$$\beta_t = \frac{\hat{p}_t}{r_t + \hat{p}_t}. \quad (35)$$

The current $s = d_k$ precision projection and the moment projection then give

$$p_t^{\text{proj}, d_k} = \hat{p}_t (1 - \beta_t), \quad p_t^{\text{mom}} = \hat{p}_t \left(1 - \frac{\beta_t}{d_k}\right). \quad (36)$$

One scalar measurement correctly removes only rank-one uncertainty, so the moment projection reduces each marginal variance by only a β_t/d_k fraction. The current rule reduces every coordinate by a β_t fraction, or d_k times as much total trace in this symmetric example. It therefore acts as a strong confidence regularizer: future gains are smaller and existing associations are overwritten less aggressively. Indeed, the current $s = 1$ projection also gives smaller posterior variance than the moment projection,

$$p_t^{\text{proj},1} = \frac{\hat{p}_t d_k r_t}{d_k r_t + \hat{p}_t} < p_t^{\text{mom}}. \quad (37)$$

Because the empirical information-scale sweep found $s = d_k$ substantially better than $s = 1$, an unretuned moment projection is likely to retain still more uncertainty, produce larger future write/erase gains, and increase interference. It may nevertheless help tasks that require rapid adaptation or frequent replacement of old associations. Posterior faithfulness and memory quality are therefore distinct hypotheses.

The process-noise scale makes a raw comparison especially unfavorable to the moment variant. Prediction injects $O(d_k \omega_t)$ total diagonal uncertainty per token, whereas a rank-one moment update removes only $O(1)$ total uncertainty. Near $\alpha_t^2 = 1$, matching the stationary effective gain of the current $s = d_k$ rule requires approximately

$$\omega_t^{\text{mom}} \simeq \frac{\omega_t^{\text{proj}}}{d_k}, \quad (38)$$

or, equivalently at the level of the stationary covariance-to-noise ratio, an observation-noise scale roughly d_k times larger. Without this recalibration, the current initialization $\omega_t \simeq \omega_{\min} + \text{softplus}(0)$ can drive the moment filter toward high uncertainty and β_t^{eff} close to one.

The appropriate comparison is therefore four-way:

1. the production Möbius projection with $s = d_k$;
2. the same projection with $s = 1$;
3. the recurrent moment projection with the same ω/r parameterization;
4. the recurrent moment projection with ω initialized and floored at approximately $1/d_k$ of the production scale, or otherwise calibrated to match the initial and stationary β_t^{eff} distribution.

The raw comparison measures the behavior of each formula under identical parameters; the gain-matched comparison isolates posterior geometry from overwrite temperature. Both should log β_t^{eff} , covariance and process-noise quantiles, effective memory half-life, and write/erase magnitude. Repeated-key and orthogonal-key synthetic tasks should precede language model pretraining, because they directly separate rapid rewriting from protection against cross-talk. The moment projection should replace the current rule only if it improves these matched comparisons enough to compensate for losing the parallel Möbius gain scan.

6.5 Anisotropic write direction and initialization

For a unit key, decompose

$$\beta_t^{\text{eff}} = \mathbf{k}_t^\top \boldsymbol{\kappa}_t, \quad \boldsymbol{\kappa}_{t,\perp} = \boldsymbol{\kappa}_t - \beta_t^{\text{eff}} \mathbf{k}_t, \quad \mathbf{k}_t^\top \boldsymbol{\kappa}_{t,\perp} = 0. \quad (39)$$

IsoKLA enforces $\boldsymbol{\kappa}_{t,\perp} = 0$: uncertainty changes the strength of a write while preserving the binding direction. DiagKLA can rotate both the write and erase directions away from the current key. This

can be useful uncertainty preconditioning, or it can write the value into unrelated key directions and create cross-talk. Greater covariance expressivity alone does not distinguish these interpretations.

The current initialization is also not an exact IsoKLA fallback. Zeroing the output map of the process-noise projection makes $\omega_{t,i}$ uniform, but

$$\hat{p}_{t,i} = \alpha_{t,i}^2 / \mu + \omega_t \quad (40)$$

is already channel dependent when $\alpha_{t,i}$ is channel dependent. After one observation, $d_k k_{t,i}^2 / r_t$ further leaves the isotropic subspace. The initialization is a useful stabilizer, but the model cannot choose an exact IsoKLA function and add anisotropy only when it is beneficial.

6.6 Decisive ablation

The cleanest next experiment holds the diagonal information tracker and effective write strength fixed while intervening only on the gain direction:

$$\kappa_t^{(\eta)} = \beta_t^{\text{eff}} \mathbf{k}_t + \eta \kappa_{t,\perp}, \quad \eta \in \{0, 1\}. \quad (41)$$

Here $\eta = 0$ removes only the off-key component and $\eta = 1$ recovers the current DiagKLA. A gain for $\eta = 0$ would implicate the off-key direction; parity would show little measurable benefit from that direction; a robust gain for $\eta = 1$ would provide direct evidence that anisotropic writes add useful capacity. The comparison should log β^{eff} , gain angle, $\mathbf{p}$ and $\mathbf{c}$ quantiles, ω/r , and cold-channel fraction on current checkpoints.

Two supporting controls are needed:

1. Compare the projected information state in Equation (15) with the sequential exact-posterior diagonal in Equation (14) on a short synthetic task, including orthogonal keys with identical elementwise squares.
2. Run a matched three-seed 220M comparison of DiagKLA, IsoKLA, and the stabilized IsoKLA variant with identical world size, data order, and parameter count. Validate the S-NIAH-2 advantage across context lengths and needle depths before spending on a new large-scale run.

A scale point above the current boundary is also warranted. Test $2d_k$ first and add $4d_k$ only if performance continues to improve. This is an overwrite-temperature diagnostic: larger s increases the post-write confidence state. A token-local change does not strengthen that token’s write, while a global change alters later current writes through their covariance prefix. The sweep can locate a better retention–plasticity operating point, but it cannot resolve the posterior inconsistency or establish that off-key anisotropy is useful.

7 Conclusion

The diagonal KDN has a genuine common-scale non-identifiability in (ω, r, μ) , but the completed noisy-MQAR experiments show that random initial anisotropy is the dominant observed instability. The appropriate design is therefore twofold: start the gain isotropically, and fix a covariance unit—preferably $r = 1$ and $\mu = 1$ —so the model learns only uncertainty ratios that affect its predictions.